

Distributed Chess Engine Using MPI

Python 3.8+

MPI mpi4py

License MIT

A distributed chess engine built with Python, leveraging the Message Passing Interface (MPI) for parallel computation of chess moves. The engine uses a minimax algorithm with alpha-beta pruning, distributed across multiple processes, and features a graphical user interface (GUI) powered by Pygame. Play chess against an AI with configurable difficulty levels, visualize moves, and explore parallel computing in action.

Table of Contents

- [Features](#)
- [Architecture](#)
- [Prerequisites](#)
- [Installation](#)
- [Usage](#)
- [Project Structure](#)
- [Configuration](#)
- [Contributing](#)
- [License](#)

Features

- **Distributed Minimax Search:** Parallelizes chess move computation using MPI, speeding up the minimax algorithm with alpha-beta pruning.
- **Interactive GUI:** Built with Pygame, offering a user-friendly chessboard with move highlights, promotion dialogs, and game status updates.
- **Configurable Difficulty:** Three levels (Easy, Medium, Hard) with adjustable search depth and time limits.
- **Advanced Evaluation:** Considers material, piece-square tables, mobility, pawn structure, king safety, and game phase.
- **Robust Error Handling:** Manages timeouts, worker failures, and invalid moves gracefully.
- **Logging:** Detailed logs for debugging and performance analysis.

Architecture

The system follows a master-worker architecture using MPI:

- **Master Process (Rank 0):** Runs the GUI (`chess_gui.py`) and coordinates task distribution (`master.py`).
- **Worker Processes (Rank > 0):** Compute minimax searches (`worker.py`) using the chess engine (`chess_engine.py`).
- **Shared Utilities:** Serialization and logging (`utils.py`) support communication and debugging.

Below is the architecture diagram (Mermaid syntax):

```
graph TD
    subgraph "Master Process (Rank 0)"
        A[chess_gui.py] --> B[GUI Thread]
        A --> C[AI Move Thread]
        C --> D[master.py]
        D --> E[Task Distribution]
        D --> F[Result Collection]
    end

    subgraph "Worker Processes (Rank > 0)"
        G[worker.py] --> H[Task Reception]
        H --> I[chess_engine.py]
        I --> J[Minimax Search]
        J --> K[Board Evaluation]
        I --> L[Result Transmission]
    end

    subgraph "Shared Components"
        M[chess_engine.py] --> N[Minimax Algorithm]
        M --> O[Evaluation Function]
        P[utils.py] --> Q[Serialization]
        P --> R[Logging]
    end

    E -- "MPI Send (FEN, Depth, Time)" --> H
    L -- "MPI Receive (Score, Move)" --> F
    Q -- "Used by" --> E
    Q -- "Used by" --> L
    R -- "Used by" --> A
    R -- "Used by" --> D
    R -- "Used by" --> G
    R -- "Used by" --> I
    B --> O
    C --> N

    style A fill:#f9f,stroke:#333,stroke-width:2px
    style D fill:#bbf,stroke:#333,stroke-width:2px
    style G fill:#bfb,stroke:#333,stroke-width:2px
    style M fill:#ffb,stroke:#333,stroke-width:2px
    style P fill:#fdd,stroke:#333,stroke-width:2px
```

PROF

Prerequisites

- **Python:** 3.8 or higher
- **MPI Implementation:** OpenMPI or MPICH
- **Python Libraries:**
 - **mpi4py:** For MPI communication

- `python-chess`: For chess logic
- `pygame`: For the GUI
- **Operating System**: Linux, macOS, or Windows (with MPI support)

Installation

1. Clone the Repository:

```
git clone https://github.com/your-username/distributed-chess-engine.git
cd distributed-chess-engine
```

2. Install MPI:

- **Ubuntu/Debian:**

```
sudo apt-get install openmpi-bin openmpi-common libopenmpi-dev
```

- **macOS:**

```
brew install openmpi
```

- **Windows:** Install WSL2 and follow Ubuntu instructions, or use MPICH.

3. Set Up a Virtual Environment (recommended):

```
python -m venv myenv
source myenv/bin/activate # On Windows: venv\Scripts\activate
```

4. Install Python Dependencies:

```
pip install mpi4py python-chess pygame
```

5. Download Chess Piece Images (optional):

- Place chess piece images (e.g., `wp.png`, `bp.png`, etc.) in a `chess_pieces` folder.
- Without images, the GUI falls back to text rendering.
- Example source: [Chess Piece Sprites](#).

Usage

1. Run the Game:

- Launch the master and worker processes using **mpirun**:

```
mpirun -n 8 --oversubscribe python chess_gui.py
```

- **-n 8**: Runs 1 master + 7 workers.
- **--oversubscribe**: Allows more processes than CPU cores (optional if cores suffice).
- **Increase Nodes**: To use more workers, increase **-n**. For example:

```
mpirun -n 16 --oversubscribe python chess_gui.py
```

- This runs 1 master + 15 workers. Optimal performance is typically with 10–35 workers, matching the number of legal moves (~35).

2. Gameplay:

- **Select Difficulty**: Choose Easy, Medium, or Hard at startup.
- **Make Moves**: Click a piece to select, then click a highlighted square to move.
- **Pawn Promotion**: Select a piece (Queen, Rook, Bishop, Knight) in the promotion dialog.
- **New Game**: Press **N** to start a new game.
- **Undo Move**: Press **U** to undo the last human and AI moves.
- **Quit**: Close the window or press **Ctrl+C** (sends STOP signal to workers).

3. Logs:

- Check **rank_X.txt** files for logs from each process (master: **rank_0.txt**, workers: **rank_1.txt**, etc.).

Project Structure

PROF

```
distributed-chess-engine/
├─ chess_gui.py      # GUI and master process entry point
├─ master.py         # Task distribution and result collection
├─ worker.py         # Worker process logic
├─ chess_engine.py   # Minimax algorithm and board evaluation
├─ utils.py          # Serialization and logging utilities
├─ chess_pieces/     # (Optional) Chess piece images
├─ rank_X.txt        # Log files per process
├─ README.md         # This file
└─ report.md         # Detailed project report
```

Configuration

Adjust AI difficulty in **chess_gui.py** under **DIFFICULTY_PROFILES**:

```
DIFFICULTY_PROFILES = {
    "Easy": {'max_depth': 3, 'time_limit': 3, 'label': 'Easy',
'use_distributed': True},
    "Medium": {'max_depth': 4, 'time_limit': 5, 'label': 'Medium',
'use_distributed': True},
    "Hard": {'max_depth': 6, 'time_limit': 10, 'label': 'Hard',
'use_distributed': True}
}
```

- `max_depth`: Maximum search depth.
- `time_limit`: Time (seconds) per move.
- `use_distributed`: Enable/disable MPI (set to `False` for local search).

Contributing

Contributions are welcome! To contribute:

1. Fork the repository.
2. Create a feature branch (`git checkout -b feature/your-feature`).
3. Commit changes (`git commit -m "Add your feature"`).
4. Push to the branch (`git push origin feature/your-feature`).
5. Open a pull request.

Please include tests and update documentation as needed.

Ideas for Improvement

- Add transposition tables for faster searches.
- Implement an opening book or neural network evaluation.
- Enhance the GUI with move history or analysis.
- Optimize load balancing for uneven subtrees.

PROF

License

This project is licensed under the [MIT License](#).

Developed as part of an academic project at **IIIT Hyderabad**, for educational and research purposes.

Repository

For future updates and changes, visit the [GitHub Repository](#).
